

Distributionally Robust Classifiers in Sentiment Analysis

Allan Li

Dept. of Computer Science
Stanford University

shilun@stanford.edu

Renee Li

Dept. of Computer Science
Stanford University

reneeli@stanford.edu

Carina Zhang

Dept. of Statistics
Stanford University

carinaz@stanford.edu

1 Introduction

Modern state-of-the-art sentiment classifiers have achieved high performance on experiments where training data and test data are from the same distribution. However, it has been shown that the classifiers do not generalize well when this assumption is broken. Recent research has developed a framework, Distributionally Robust Optimization (DRO), to offset the negative influence of shifting data distribution and make the underlying classification models perform better when there is a change in distribution (Duchi and Namkoong, 2018). While heuristic methodologies have shown progress and were tested on synthetic datasets and computer vision datasets, we want to extend DRO to Natural Language Processing tasks, such sentiment analysis.

In this project, we investigate how sentiment classification models integrated with DRO will behave under distributional shifts. To narrow down the scope, we will consider one form of distributional shift and understand how it affects our models.

Our major contributions are:

- successfully applying the DRO framework to develop machine learning models (specifically sentiment analysis within the field of NLP) when the test set’s population and the training set’s population do not overlap but are similar in nature;
- confirming through our experiments that our DRO model does improve performance on our test set with distributional shift from the training set;
- developing key observation that while there is no obvious pattern with small radii, the performance of our DRO model converges to

the performance of baseline model as the radius of the L_p ball gets very large; thus, the best performing radius is usually somewhere in the middle.

2 Related Works

Recent research attempted to develop a model that can perform well across all unforeseen test distributions. While Duchi et. al. were able to develop and analyze a DRO framework that learns a model that provides good performance against perturbations to the data-generating distribution (Duchi and Namkoong, 2018), we want to explore whether a model can be developed to be trained and tested on NLP datasets as well to show the increasing generality of the DRO framework.

Previous work in domain adaptation leads to development of models that receive data from one domain and are tested on a specified target. For example, there have been NLP-specific implementations of DRO (Oren et al., 2019; Anonymous, 2020). Some researchers describe a framework that deals with a specific distributional shift, *sub-population shift*, in which the test distribution is a subpopulation of the training distribution. They have achieved better results than generic DRO and have shown that topics are an effective way to encode prior information about test distributions (Oren et al., 2019). A stochastic optimizer is also introduced to a DRO framework that scales to large models and datasets, while improving robust accuracy at only a small cost in average accuracy (Anonymous, 2020).

Building upon these implementations, we will explore another type of distributional shift - when the training set’s population and test set’s population does not overlap but are similar in nature and have the same set of output labels.

3 Movie Review Datasets

We choose two movie review datasets, IMDB and Rotten Tomatoes, as our training set and test set. Since Rotten Tomatoes reviews and IMDB are similar data sets and generate either a "positive" or a "negative" sentiment output, Rotten Tomatoes can be considered as a shifted IMDB test set.

3.1 Training Set: IMDB Reviews

We use Stanford AI Lab's Large Movie Review Dataset (Maas et al., 2011) as our training set. The dataset was collected from IMDB and contains roughly 50,000 movie reviews, in which 25,000 are positive and 25,000 are negative. The core dataset contains 50,000 reviews in total, we are using the 25,000 reviews from the training set. No more than 30 reviews are selected from the reviews pool of any single movie because reviews for the same movie tend to have correlated ratings.

3.2 Test Set: Rotten Tomatoes Reviews

We acquire labeled Rotten Tomatoes movie review dataset from the Stanford Treebank Project (Socher et al., 2013). Since we want to see how DRO performs on shifted a dataset, it seems most fitting for us to employ another large dataset of movie reviews as our test set. This dataset is split into train, dev, and test sets containing 8,544, 1,101, and 2,210 reviews respectively. We have used the pre-split test set as our test set.

4 Model and Algorithm

This section describes our model and algorithm in detail. The baseline and oracle models and experiments are described in Section 5.

4.1 Advanced Model Description

We use PyTorch pre-trained BERT-base-based model (referred to as BERT later) to produce word embeddings and use 2 layer bi-LSTM and linear classifier as our underlying classifier for sentiment classification. The pre-trained BERT is illustrated in Figure 1. We then implement DROs on top of these classifiers and record the shifted test accuracy. Alternatives that we can potentially explore later would include SVM, Softmax, and Convolutional Neural Networks as classifiers.

In order to help our model better understand the reviews, we propose a model architecture (see Figure 2). As a comparison to just using DROs,

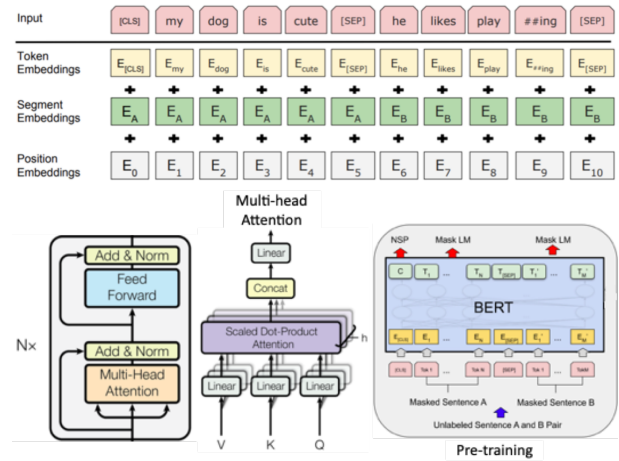


Figure 1: BERT's key technical innovation is applying the bidirectional training of Transformer, a popular attention model, to language modelling, as shown in this example.

we also did a controlled experiment where we trained BERT on IMDB dataset and tested on Rotten Tomatoes just using BERT. The architecture can be separated into 5 stages.



Figure 2: High level model architecture.

- The first stage is the data pre-processing step. We have made some improvements to the pre-processing procedures from the one mentioned in our baseline model. We will discuss this step in more detail in Section 5.
- In the second stage, we pass in word pieces from the dataset of movie reviews to the state-of-the-art language model for NLP – BERT.
- At stage three, we pass the result we get from BERT into 2 bi-LSTM layers. While we are at the latter layer, we concatenated its forward and backward together.

- d. We project the result onto a L_p -ball or a simplex for the robust model.
- e. We apply dropout and then give the output to a linear classifier.
- f. Finally, we use a sigmoid function to wrap up this model. Results and model comparison will be discussed in Section 4.

4.2 Formulation of DRO

Performance of machine learning models degrades significantly on test sets that are different from what the model was trained on because of its reliance on a priori fixed target distribution (Oren et al., 2019; Duchi and Namkoong, 2018; Anonymus, 2020). To mitigate these challenges, we use a loss minimization framework that is explicitly robust to local changes in the data distribution. Concretely, let $\Theta \in R^d$ be the parameter space, P_0 be the data distribution, and $\ell(\theta; X)$ as the total loss function. When DRO models make predictions, rather than minimizing the average loss, $E[\ell(\theta; X)]$, we study the distributionally robust stochastic optimization problem,

$$\min_{\theta \in \Theta} \{R_f(\theta; P_0) := \sup_{Q \ll P_0} \{E_Q[\ell(\theta; X)]\} \quad (1)$$

We call (1) the worst case risk, which captures how a model can performance in a worse-case situation. Optimizing worst-case situation is the crux of model robustness and generalizability. Since the worst-case risk upweights regions of X with high losses, it consequently optimizes performance on the tails. The model parameters that achieves the worst-case risk will be the parameters of the DRO model (Duchi and Namkoong, 2018).

4.3 Model workflow with examples

As mentioned in our proposal, we use two datasets, IMDB movie reviews dataset and Rotten Tomatoes movie reviews dataset, as training sets and test sets.

The pre-processing step is explained in details in Section 5 in our demonstration of our experiment.

We tokenize the sentences to feed into BERT. Since BERT has a constraint of a maximum token length of 512, we only take the first 512 tokens of a sentence to represent that sentence. After data processing, each movie review can be represented as a list of BERT token indexes of length d , where $d = 512$ in our context.

Since neural networks spill out an unbounded intermediate output, y' , we need to apply sigmoid function $\hat{y} = \frac{1}{1+e^{-y'}}$ to bound have a probability normalized between (0, 1).

Since sentiment classification is a binary classification problem, we investigate distributional shifts via a binary classification experiment using the binary cross entropy with sigmoid, defined as $\ell(\theta; x, y) = \frac{1}{N} \sum_{i=1}^N y_i \cdot \log(\hat{y}_i) + (1 - y_i) \cdot \log(1 - \hat{y}_i)$.

As a **concrete example**, our training data from IMDB will look like this: "If you like original gut wrenching laughter you will like this movie. If you are young or old then you will love this movie, hell even my mom liked it. Great Camp!!!" and is labeled as "positive". Our test data (after pre-processing) from Rotten Tomatoes for the model would look like this: "Take Care of My Cat offers a refreshingly different slice of Asian cinema ." The correct prediction should be "1" because it is a sentence with a "positive" sentiment.

Last but not least, since PyTorch's LSTM requires that the batch is sorted by decreasing token length, we complete the sorting of each batch while collecting the batches so that when each batch is passed into the biLSTM layer, it has already been sorted in the order we need.

The workflow with example inputs/outputs and processes is visualized in 3 below.

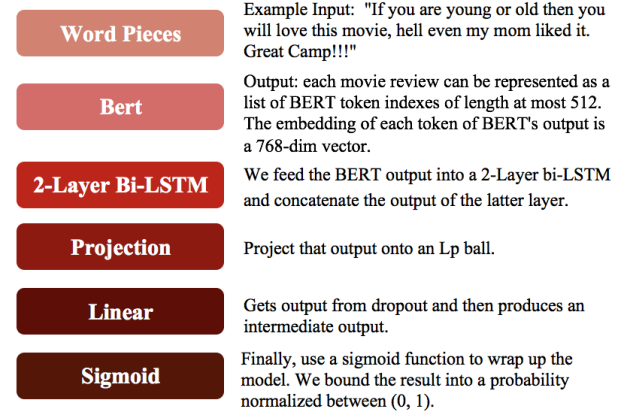


Figure 3: Model Architecture with example inputs/outputs and process.

4.4 Potential challenges

1. Since we are dealing with truly unforeseen test data in our experiment, how much our model can improve from the baseline is not guaranteed.

2. It is hard to mathematically define the distributional shift explored in our experiment. It is obvious that IMDB reviews and Rotten Tomatoes reviews are similar datasets for a human being; however, it is hard to quantify the magnitude of the shift, which makes the results of our project hard to generalize beyond movie reviews.

5 Experiment

In this section, we first introduce the evaluation metrics we use for the models. Then we introduce the methods we adopt to preprocess the text information. Lastly, we describe the experiments we conducted in our project.

5.1 Evaluation Design and Metrics

1. Independent Variables: Test Accuracy
2. Dependent Variables:
 - (a) Base classifier: the model used to accomplish sentiment classification task.
 - (b) Test set distributional shift: the distribution of data that we expose the models to in testing.

5.2 Data pre-processing

For the training set, we use the labeled IMDB train set for this project. In the labeled training set, a negative review has a score ≤ 4 out of 10, and a positive review has a score ≥ 7 out of 10. Reviews with neutral ratings (with scores of 5 or 6) are not included in the training sets.

For the test set, since our task is binary classification, review scores have a range from 0 to 4. A negative review has a score < 2 , and a positive review has a score > 2 . Reviews with score equal to 2 are not included in our dataset.

We filter the dataset as described above and output the sentences with sentiment labels, where 1 indicates positive sentiment and 0 indicates negative sentiment. We output the parsed dataset to a .csv file, which speeds up the process of data loading during training.

As a **concrete example**, our input data (after pre-processing) for the model would look like this: "Take Care of My Cat offers a refreshingly different slice of Asian cinema." And the output would be "1" because it is a sentence with a "positive" sentiment.

5.3 Baseline Experiment

We implement GloVe model, a relatively naive sentiment classifier, as our baseline model (Pennington et al., 2014). We trained our GloVe on IMDB data set and tested on Rotten Tomatoes test set. We use test accuracy obtained by GloVe without applying DRO as our baseline accuracy, which is 77.7%.

5.4 Oracle Experiment

We choose the best performing model in terms of in-distribution test accuracy on the labeled Rotten Tomatoes reviews dataset, ALBERT, which is a lighter version of BERT model with fewer parameters with regularization. This model achieved an in-distribution test accuracy of 97.1% (Lan et al., 2019).

5.5 Advanced Method Experiment

1. Encoding: We use contextualized word embedding produced by BERT from word pieces to encode the text inputs, which are movie reviews stored as a string. Since fine-tuning the model requires too much time, we freeze the pre-trained BERT to get a fixed representation of words to get the vector presentation to feed into DRO.
2. Training: We train using a subset of the different types of classifiers mentioned above on the IMDB data set after the pre-process procedures as mentioned above.
3. DRO Implementation: We have attempted the implementation of several DROs to our classifiers. In particular, we have tried projecting on an L_p -ball where $p = 1, 2$ or 4 and a simplex.
4. Testing: We test our DRO-integrated models (i.e. robust models) on the Rotten Tomatoes test set and compare the results between in-distribution test accuracy and shifted test accuracy and evaluate using the metrics explained below.

6 Results and Discussion

We implemented two non-robust classification models for comparison purposes and two robust models with L_1 , L_2 , L_4 and Simplex with give different radii, $R = 1, 2, 5, 10, 15$. We report test accuracy for each model.

Table 1: Results of non-robust model performance

Train set	Test Set	Test Accuracy
IMDB	IMDB	83.60%
IMDB	R. Tomatoes	73.64%

6.1 Non-Robust Model Performance

We implemented our classification model by the procedure we described in the previous sections without implementing DRO. Thus, this classification model based on pre-trained BERT and linear classifier can be viewed as a baseline of how classification models perform under distributional shifts. The results of the model is shown in Table-1. We can see from the table that there exists a significant drop in performance (around 10%) when the test set comes from a different population.

6.2 DRO Models Performance

We added the DRO implementation described in previous sections with different model configurations, namely Simplex, L_1 , L_2 , L_4 projections with radii $R = 1, 2, 5, 10, 15$. We report test accuracy without distributional shift (referred as in-distribution test accuracy, namely trained on IMDB and tested on IMDB) and test accuracy with distributional shift (referred as out-of-distribution test accuracy, namely trained on IMDB and tested on Rotten Tomatoes) for each model. The overall performance is reflected in Table-2 to Table-5 and plotted in Figure-4.




Figure 4: Models' out-of-distribution accuracy with difference configurations

6.3 Analysis and Interpretation of Experimental Results

Comparing results from our DRO models to that of the non-robust model, our observations align with

Table 2: DRO Model with L_1 Projection

Radius	In-distribution Test Accuracy	Out-of-distribution Test Accuracy
1	80.65%	74.72%
2	75.10%	69.34%
5	84.55%	80.62%
10	84.25%	78.10%
15	85.00%	80.04%

Table 3: DRO Model with L_2 Projection

Radius	In-distribution Test Accuracy	Out-of-distribution Test Accuracy
1	84.90%	81.02%
2	84.85%	81.13%
5	83.45%	75.75%
10	85.60%	82.45%
15	84.90%	78.90%

Table 4: DRO Model with L_4 Projection

Radius	In-distribution Test Accuracy	Out-of-distribution Test Accuracy
1	84.90%	80.10%
2	85.50%	80.39%
5	85.00%	80.73%
10	84.95%	80.16%
15	85.10%	79.70%

Table 5: DRO Model with Simplex Projection

Radius	In-distribution Test Accuracy	Out-of-distribution Test Accuracy
1	84.45%	78.04%
2	83.60%	81.02%
5	84.70%	78.21%
10	84.35%	78.84%
15	84.20%	76.38%

the assumption that DRO performs better under distributional shifts. When exposed to shifted data, the certified test accuracy of our model decreases from above 90% to 77%. Thus, this result has shown that when models are exposed to distributional shift (in our context, the population changes from IMDB and Rotten Tomatoes). We can infer that the current state-of-the-art sentiment classifier is not robust under distributional shifts.

One potential cause of the drop in performance is that empirically, there exists an inherent trade-off between robustness and accuracy. Increasing robustness requires more relaxation and thus may result in lower certified test accuracy for a particular set of perturbations. Since the three models studied in our paper are tuned to a specific fixed, known family of distortions, the generalizability of these models decreases.

An observation that aligns with this potential cause of drop in performance is the observation that test accuracy and radius does not have a linear relationship. Intuitively, the smaller R is, the more "relaxed" the optimization problem is. Larger R might lead to convergence to in-distribution results. However, since DRO models are not tuned to fit specifically to this specific distributional shift, there is a robustness-accuracy trade off. Thus, a radius value in the middle usually produce the best results.

We identified two potential sources of error. First, we see that sometimes our model's performance dropped below the baseline. Because of the limitation in time and computational power, we did not have time to run every experiment several times and record the average result. Thus, randomness causes the variance in model performance. Second, the challenge lies in the fact that the effect of projection on the original input in a natural language setting is more ambiguous than images. Since it is hard to mathematically define our distributional shift, it is also hard to precisely define the projection's effect on sentences. Some of the projection methods, such as L_1 , is not suitable in this setting. However, we included this model configuration for the sake of completeness.

7 Conclusion

In this work, we propose sentiment classification models integrated with DRO to improve model performance on datasets with distributional shifts. To narrow down the scope, we consider one form

of distributional shift (from IMDB dataset to Rotten Tomatoes dataset). We evaluate our model's performance based on test accuracy. We conduct experiments and show that the model's performance does not have a linear relationship with the radius of the L_p balls we project onto; instead, the performance gradually converges to baseline performance with the increase of radius. We found that out of Simplex, L_1 , L_2 , and L_4 projections, L_4 had the best overall performance.

The major contributions of this work include applying the DRO framework to machine learning models (specifically sentiment analysis within the field of NLP) when the test set's population and the training set's population do not overlap but are similar in nature. Meanwhile, we have confirmed through our experiments that our DRO model does improve performance on our test set with distributional shift from the training set.

7.1 Future Work

While we successfully accomplished our task detailed in the Introduction section through this project, there are potential directions we want to explore and questions we want to answer in the future. Overall, we want to look into whether we could add more diversity to our distributional shifts beyond a total population shift. More specifically, we can experiment with mixing up IMDB and Rotten Tomatoes datasets with different ratios to have different distributions to test on. We are also interested in further exploring with datasets for which we can better quantify the underlying distributional shift.

Acknowledgement

We would like to acknowledge the help of Professor Percy Liang, Professor Dorsa Sadigh, and Professor John Duchi for brainstorming the project idea, Dr. Hongseok Namkoong for the initial codebase. We would also like to thank our amazing teaching assistant Haoshen Hong for providing helpful suggestions and feedback as well as effective solutions to tackle the problem.

Authorship

Allan Li, Renee Li, and Carina Zhang designed and conducted the research; Allan Li processed the dataset, implemented the final model and collected final results; Renee Li and Carina wrote the paper. All authors contributed equally.

References

- Anonymous. 2020. Distributionally robust neural networks. In *Submitted to International Conference on Learning Representations*.
- John Duchi and Hongseok Namkoong. 2018. Learning models with uniform performance via distributionally robust optimization. *arXiv preprint arXiv:1810.08750*.
- Zhenzhong Lan, Mingda Chen, Sebastian Goodman, Kevin Gimpel, Piyush Sharma, and Radu Soricut. 2019. Albert: A lite bert for self-supervised learning of language representations. *arXiv preprint arXiv:1909.11942*.
- Andrew L. Maas, Raymond E. Daly, Peter T. Pham, Dan Huang, Andrew Y. Ng, and Christopher Potts. 2011. [Learning word vectors for sentiment analysis](#). In *Proceedings of the 49th Annual Meeting of the Association for Computational Linguistics: Human Language Technologies - Volume 1*, HLT '11, pages 142–150, Stroudsburg, PA, USA. Association for Computational Linguistics.
- Yonatan Oren, Shiori Sagawa, Tatsunori B Hashimoto, and Percy Liang. 2019. Distributionally robust language modeling. *arXiv preprint arXiv:1909.02060*.
- Jeffrey Pennington, Richard Socher, and Christopher Manning. 2014. Glove: Global vectors for word representation. In *Proceedings of the 2014 conference on empirical methods in natural language processing (EMNLP)*, pages 1532–1543.
- Richard Socher, Alex Perelygin, Jean Wu, Jason Chuang, Christopher Manning, Andrew Ng, and Christopher Potts. 2013. Parsing with compositional vector grammars. In *EMNLP*.

Supplementary Material

Link to github: <https://github.com/Shilun-Allan-Li/CS221-Fall2019>